
Extreme Q-Learning: MaxEnt RL without Entropy

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Modern Deep Reinforcement Learning (RL) algorithms require estimates of the
2 maximal Q-value, which are difficult to compute in continuous domains with an
3 infinite number of possible actions. In this work, we introduce a new update rule
4 for online and offline RL which directly models the maximal value using Extreme
5 Value Theory (EVT) inspired by Economics. By doing so, we avoid computing
6 Q-values using out-of-distribution actions which is often a substantial source of
7 error. Our key insight is to introduce an objective that directly estimates the optimal
8 soft-value functions (LogSumExp) in the maximum entropy (MaxEnt) RL setting
9 without needing to sample from a policy. Using EVT, we derive our *Extreme*
10 *Q-Learning* framework and consequently online and, for the first time, offline
11 MaxEnt Q-learning algorithms, that do not explicitly require access to a policy or
12 its entropy. Finally, our method obtains state-of-the-art results in the Offline D4RL
13 benchmark strongly outperforming prior works by *10-20 points* while obtaining
14 moderate improvements over SAC and TD3 on online DM Control tasks.

15 1 Introduction

16 Modern Deep Reinforcement Learning (RL) algorithms have shown broad success in challenging
17 control [14, 33] and game-playing domains [25]. While tabular Q-iteration or value-iteration methods
18 are well understood, state of the art RL algorithms often make theoretical compromises in order to
19 deal with deep networks, high dimensional state spaces, and continuous action spaces. In particular,
20 standard Q-learning algorithms require computing the max or soft-max over the Q-function in order
21 to fit the Bellman equations. Yet, almost all current off-policy RL algorithms for continuous control
22 only indirectly estimate the Q-value of the next state with separate policy networks. Consequently,
23 these methods only estimate the Q-function of the current policy, instead of the optimal Q^* , and rely
24 on policy improvement via an actor. Moreover, actor-critic approaches on their own have shown to
25 be catastrophic in the offline settings where actions sampled from a policy are consistently out-of-
26 distribution [19, 10]. As such, computing $\max Q$ for Bellman targets remains a core issue in deep RL.

27 One popular approach is to train Maximum Entropy (MaxEnt) policies, in hopes that they are more
28 robust to modeling and estimation errors [42]. However, the Bellman backup \mathcal{B}^* used in MaxEnt
29 RL algorithms still requires computing the log-partition function over Q-values, which is usually
30 intractable in high-dimensional action spaces. Instead, current methods like SAC [14] rely on auxiliary
31 policy networks, and as a result do not estimate \mathcal{B}^* , the optimal Bellman backup. Our key insight
32 is to apply extreme value analysis used in branches of Finance and Economics to Reinforcement
33 Learning. Ultimately, this will allow us to directly model the LogSumExp over Q-functions in the
34 MaxEnt Framework.

35 Intuitively, reward or utility-seeking agents will consider the maximum of the set of possible future
36 returns. The Extreme Value Theorem (EVT) tells us that maximal values drawn from any exponential
37 tailed distribution follows the Generalized Extreme Value (GEV) Type-1 distribution, also referred to
38 as the Gumbel Distribution \mathcal{G} . The Gumbel distribution is thus a prime candidate for modeling errors

in Q-functions. In fact, McFadden’s 2000 nobel-prize winning work in Economics on discrete choice models [24] showed that soft-optimal utility functions with logit (or softmax) choice probabilities naturally arise when utilities are assumed to have Gumbel-distributed errors. This was subsequently generalized to stochastic MDPs by Rust [32] in 1984. Nevertheless, these intriguing results have remained largely unknown in the RL community. By introducing a novel loss optimization framework, we bring them into the world of modern deep RL.

Empirically, we find that even modern deep RL approaches, for which errors are typically assumed to be Gaussian, exhibit errors that better approximate the Gumbel Distribution, see Figure 1. By assuming errors to be Gumbel distributed, we obtain Gumbel Regression, a consistent estimator over log-partition functions even in continuous spaces. Furthermore, making this assumption about Q -values lets us derive a new Bellman loss objective that directly solves for the optimal MaxEnt Bellman operator \mathcal{B}^* , instead of the operator under the current policy \mathcal{B}^π . As soft optimality gracefully emerges from our framework, we can run MaxEnt RL independently of the policy. In the online setting, we avoid using a policy network to explicitly compute entropies. In the offline setting, we completely avoid sampling from learned policy networks, minimizing the aforementioned extrapolation error. Our resulting algorithms surpass state-of-the-art (SOTA) while being practically simpler.

In this paper we outline the theoretical motivation for using Gumbel distributions in reinforcement learning, and show how it can be used to derive practical online and offline MaxEnt RL algorithms. Concretely, our contributions are as follows:

- We motivate Gumbel Regression and show it allows calculation of the log-partition function (LogSumExp) in continuous spaces. We apply it to MDPs to present a novel loss objective for RL using maximum-likelihood estimation.
- Our formulation extends soft-Q learning to offline RL as well as continuous action spaces without the need of policy entropies. It presents a way to directly calculate the optimal soft-values V^* and soft-Bellman updates \mathcal{B}^* using SGD, which are usually intractable in continuous settings.
- We provide the missing theoretical link between soft and conservative Q-learning, showing how these formulations can be made equivalent. We also show how Max-Ent RL emerges naturally from vanilla RL as a conservatism in our framework.
- Finally, we empirically demonstrate state-of-art results in Offline RL, improving over prior methods by a large margin on the D4RL benchmark, and performing moderately better than SAC and TD3 in Online RL, while theoretically avoiding actor-critic formulations.

2 Preliminaries

In this section we introduce Maximum Entropy (MaxEnt) RL and Extreme Value Theory (EVT), which we use to motivate our framework to estimate extremal values in RL.

We consider an infinite-horizon Markov decision process (MDP), defined by the tuple $(\mathcal{S}, \mathcal{A}, \mathcal{P}, r, \gamma)$, where \mathcal{S}, \mathcal{A} represent state and action spaces, $\mathcal{P}(s'|s, a)$ represents the environment dynamics, $r(s, a)$ represents the reward function, and $\gamma \in (0, 1)$ represents the discount factor. In the offline RL setting, we are given a dataset $\mathcal{D} = (s, a, r, s')$ of tuples sampled from trajectories under a behavior policy $\pi_{\mathcal{D}}$ without any additional environment interactions. We use $\rho_\pi(s)$ to denote the distribution of states that a policy $\pi(a|s)$ generates. In the MaxEnt framework, an MDP with entropy-regularization is referred to as a soft-MDP [4] and we often use this notation.

2.1 Maximum Entropy RL

Standard RL seeks to learn a policy that maximizes the expected sum of (discounted) rewards $\mathbb{E}_\pi[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)]$, for (s_t, a_t) drawn at timestep t from the trajectory distribution that π generates. We consider a generalized version of Maximum Entropy RL that augments the standard reward objective with the KL-divergence between the policy and a reference distribution μ : $\mathbb{E}_\pi[\sum_{t=0}^{\infty} \gamma^t (r(s_t, a_t) - \beta \log \frac{\pi(a_t|s_t)}{\mu(a_t|s_t)})]$, where β is the regularization strength. When μ is uniform \mathcal{U} , this becomes the standard MaxEnt objective used in online RL up to a constant. In the offline RL setting, we choose μ to be the behavior policy $\pi_{\mathcal{D}}$ that generated the fixed dataset \mathcal{D} . Consequently,

89 this objective enforces a conservative KL-constraint on the learned policy, keeping it close to the
90 behavior policy [28, 14].

91 In MaxEnt RL, the soft-Bellman operator $\mathcal{B}^* : \mathbb{R}^{\mathcal{S} \times \mathcal{A}} \rightarrow \mathbb{R}^{\mathcal{S} \times \mathcal{A}}$ is defined as $(\mathcal{B}^*Q)(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) +$
92 $\gamma \mathbb{E}_{\mathbf{s}' \sim \mathcal{P}(\cdot|\mathbf{s}, \mathbf{a})} V^*(\mathbf{s}')$ where Q is the soft-Q function and V^* is the optimal soft-value satisfying:

$$V^*(\mathbf{s}) = \beta \log \sum_{\mathbf{a}} \mu(\mathbf{a}|\mathbf{s}) \exp(Q(\mathbf{s}, \mathbf{a})/\beta) := \mathbb{L}_{\mathbf{a} \sim \mu(\cdot|\mathbf{s})}^{\beta} [Q(\mathbf{s}, \mathbf{a})], \quad (1)$$

93 where we denote the log-sum-exp (LSE) using an operator \mathbb{L}^{β} for succinctness¹. The soft-Bellman
94 operator has a unique contraction Q^* [14] given by the soft-bellman equation: $Q^* = \mathcal{B}^*Q^*$ and the
95 optimal policy satisfies [13]:

$$\pi^*(\mathbf{a}|\mathbf{s}) = \mu(\mathbf{a}|\mathbf{s}) \exp((Q^*(\mathbf{s}, \mathbf{a}) - V^*(\mathbf{s}))/\beta). \quad (2)$$

96 Instead of estimating soft-values for a policy $V^{\pi}(\mathbf{s}) = \mathbb{E}_{\mathbf{a} \sim \pi(\cdot|\mathbf{s})} [Q(\mathbf{s}, \mathbf{a}) - \beta \log \frac{\pi(\mathbf{a}|\mathbf{s})}{\mu(\mathbf{a}|\mathbf{s})}]$, our
97 approach will seek to directly fit the optimal soft-values V^* , i.e. the log-sum-exp (LSE) of Q values.

98 2.2 Extreme Value Theorem

99 The Fisher-Tippett or extreme value theorem tells us that the maximum of i.i.d. samples from exponen-
100 tially tailed distributions will asymptotically converge to the Gumbel distribution $\mathcal{G}(\mu, \beta)$, which has
101 pdf $p(x) = \exp(-(z + e^{-z}))$ where $z = (x - \mu)/\beta$ with location parameter μ and scale parameter β .

102 **Theorem 1** (Extreme Value Theorem (EVT) [26, 8]). *For i.i.d. random variables $X_1, \dots, X_n \sim f_X$,
103 with exponential tails, $\lim_{n \rightarrow \infty} \max_i(X_i)$ follows the Gumbel (GEV-I) distribution. Furthermore, \mathcal{G}
104 is max-stable, i.e. if $X_i \sim \mathcal{G}$, then $\max_i(X_i) \sim \mathcal{G}$ holds.*

105 This result is similar to the Central Limit Theorem (CLT), which states that means of i.i.d. errors
106 approach the normal distribution. Thus, under a chain of max operations, any i.i.d. exponential tailed
107 errors² will tend to become Gumbel distributed and stay as such. EVT will ultimately suggest us to
108 characterize nested errors in Q-learning as following a Gumbel distribution. In particular, the Gumbel
109 distribution \mathcal{G} exhibits unique properties we will exploit.

110 One intriguing consequence of the Gumbel’s max-stability is its ability to convert the maximum
111 over a discrete set into a softmax. This is known as the *Gumbel-Max Trick* [29, 15]. Concretely for
112 i.i.d. $\epsilon_i \sim \mathcal{G}(0, \beta)$ added to a set $\{x_1, \dots, x_n\} \in \mathbb{R}$, $\max_i(x_i + \epsilon_i) \sim \mathcal{G}(\beta \log \sum_i \exp(x_i/\beta), \beta)$,
113 and $\arg\max(x_i + \epsilon_i) \sim \text{softmax}(x_i/\beta)$. Furthermore, the Max-trick is unique to the Gumbel [23].

114 These properties lead into the McFadden-Rust model [24, 32] of MDPs as we state below.

115 **McFadden-Rust model:** An MDP following the standard Bellman equations with stochasticity in the
116 rewards due to unobserved state variables will satisfy the soft-Bellman equations over the observed
117 state with actual rewards $\bar{r}(\mathbf{s}, \mathbf{a})$, given two conditions:

- 118 1. Additive separability (AS): observed rewards have additive i.i.d. Gumbel noise, i.e.
119 $r(\mathbf{s}, \mathbf{a}) = \bar{r}(\mathbf{s}, \mathbf{a}) + \epsilon(\mathbf{s}, \mathbf{a})$, with actual rewards $\bar{r}(\mathbf{s}, \mathbf{a})$ and i.i.d. noise $\epsilon(\mathbf{s}, \mathbf{a}) \sim \mathcal{G}(0, \beta)$.
- 120 2. Conditional Independence (CI): the noise $\epsilon(\mathbf{s}, \mathbf{a})$ in a given state-action pair is conditionally
121 independent of that in any other state-action pair.

122 Moreover, the converse also holds: Any MDP satisfying the Bellman equations and following a
123 softmax policy, necessarily has any i.i.d. noise in the rewards with *AS + CI* conditions be Gumbel
124 distributed.

125 These results were first shown to hold in discrete choice theory by McFadden [24], with the *AS + CI*
126 conditions derived by Rust [32] for discrete MDPs. We formalize these results in Appendix A and
127 give succinct proofs using the developed properties of the Gumbel distribution. These results enable
128 the view of a soft-MDP as an MDP with hidden i.i.d. Gumbel noise in the rewards.

129 Notably, this result gives a different interpretation of a soft-MDP than entropy regularization to allow
130 us to recover the soft-Bellman equations.

¹In continuous action spaces, the sum over actions is replaced with an integral over the distribution μ .

²Bounded random variables are sub-Gaussian [41] which have exponential tails.

3 Extreme Q-Learning

In this section, we motivate our Extreme Q-learning framework, which directly models the soft-optimal values V^* , and show it naturally extends soft-Q learning. Notably, we use the Gumbel distribution to derive a new optimization framework for RL via maximum-likelihood estimation and apply it to both online and offline settings.

3.1 Gumbel Error Model

Although assuming Gumbel errors in MDPs leads to intriguing properties, it is not obvious why the errors might be distributed as such. First, we empirically investigate the distribution of Bellman errors by computing them over the course of training. Specifically, we compute $r(s, a) - \gamma Q(s', \pi(s')) - Q(s, a)$ for samples (s, a, s') from the replay-buffer using a single Q -function from SAC [14] (See Appendix E for more details). In Figure 1, we find the errors to be skewed and better fit by a Gumbel distribution. We explain this using EVT.




Figure 1: Bellman errors from SAC on Cheetah-Run [35]. The Gumbel distribution better captures the skew versus the Gaussian.

Consider fitting Q -functions by learning an unbiased function approximator \hat{Q} to solve the Bellman equation. We will assume we have access to M such function approximators, each of which are assumed to be independent e.g. parallel runs of a model over an experiment. We can see approximate Q -iteration as performing:

$$\hat{Q}_t(s, a) = \bar{Q}_t(s, a) + \epsilon_t(s, a), \quad (3)$$

where $\mathbb{E}[\hat{Q}] = \bar{Q}_t$ is the expected value of our prediction \hat{Q}_t for an intended target \bar{Q}_t over our estimators, and ϵ_t is the (zero-centered) error in our estimate. Here, we assume the error ϵ_t comes from the same underlying distribution for each of our estimators, and thus are i.i.d. random variables with a zero-mean. Now, consider the bootstrapped estimate using one of our M estimators chosen randomly:

$$\hat{B}^* \hat{Q}_t(s, a) = r(s, a) + \gamma \max_{a'} \hat{Q}_t(s', a') = r(s, a) + \gamma \max_{a'} (\bar{Q}_t(s', a') + \epsilon_t(s', a')). \quad (4)$$

We now examine what happens after a subsequent update. At time $t+1$, suppose that we fit a fresh set of M independent functional approximators \hat{Q}_{t+1} with the target $\hat{B}^* \hat{Q}_t$, introducing a new unbiased error ϵ_{t+1} . Then, for $\bar{Q}_{t+1} = \mathbb{E}[\hat{Q}_{t+1}]$ it holds that

$$\bar{Q}_{t+1}(s, a) = r(s, a) + \gamma \mathbb{E}_{s'|s, a} [\mathbb{E}_{\epsilon_t} [\max_{a'} (\bar{Q}_t(s', a') + \epsilon_t(s', a'))]]. \quad (5)$$

As \bar{Q}_{t+1} is an expectation over both the dynamics and the functional errors, it accounts for all uncertainty (here $\mathbb{E}[\epsilon_{t+1}] = 0$). But, the i.i.d. error ϵ_t remains and will be propagated through the Bellman equations and its chain of max operations. Due to Lemma 1, ϵ_t will become Gumbel distributed in the limit of t , and remain so due to the Gumbel distribution's max-stability.³

This highlights a fundamental issue with approximation-based RL algorithms that minimize the Mean-Squared Error (MSE) in the Bellman Equation: they implicitly assume, via maximum likelihood estimation, that errors are Gaussian. In Appendix A, we further study the propagation of errors using the McFadden-Rust MDP model, and use it to develop a simplified Gumbel Error Model (GEM) for errors under functional approximation.

In practice, the Gumbel nature of the errors may be weakened as, unlike in our model, estimators between timesteps share parameters and errors will be correlated across states and actions.

3.2 Gumbel Regression

The goal of our work is to directly model the log-partition function (LogSumExp) over $Q(s, a)$ to avoid all of the aforementioned issues with taking a max in the function approximation domain.

³Same holds for soft-MDPs as log-sum-exp can be expanded as a max over i.i.d. Gumbel random vars.




Figure 2: **Left:** The pdf of the Gumbel distribution with $\mu = 0$ and different values of β . **Center:** Our Gumbel loss for different values of β . **Right:** Gumbel regression applied to a two-dimensional random variable for different values of β . The smaller the value of β , the more the regression fits the extrema.

175 In this section we derive an objective function that models the LogSumExp by simply assuming
 176 errors to follow a gumbel distribution.

177 Consider estimating a parameter h for a random variable X using samples x_i from a dataset \mathcal{D} ,
 178 which have Gumbel distributed noise, i.e. $x_i = h + \epsilon_i$ where $\epsilon_i \sim -\mathcal{G}(0, \beta)$. Then, the average
 179 log-likelihood of the dataset \mathcal{D} as a function of h is given as:

$$\mathbb{E}_{x_i \sim \mathcal{D}} [\log p(x_i)] = \mathbb{E}_{x_i \sim \mathcal{D}} \left[-e^{((x_i - h)/\beta)} + (x_i - h)/\beta \right] \quad (6)$$

180 Maximizing the log-likelihood yields the following convex minimization objective in h ,

$$\mathcal{L}(h) = \mathbb{E}_{x_i \sim \mathcal{D}} \left[e^{(x_i - h)/\beta} - (x_i - h)/\beta - 1 \right] \quad (7)$$

181 which forms our objective function $\mathcal{L}(\cdot)$ ⁴ for Gumbel Regression. β is fixed as a hyper-parameter,
 182 and we show its affect on the loss in Figure 2. Critically, the minima of this objective under a fixed β
 183 is given by $h = \beta \log \mathbb{E}_{x_i \sim \mathcal{D}} [e^{x_i/\beta}]$, which resembles the LogSumExp with the summation replaced
 184 with an (empirical) expectation.

185 In fact, this solution is the the same as the operator $\mathbb{L}_\mu^\beta(X)$ defined for MaxEnt in Section 2.1 with x_i
 186 sampled from μ . In Figure 2, we show plots of Gumbel Regression on a simple dataset with different
 187 values of β . As this objective recovers $\mathbb{L}^\beta(X)$, we next use it to model soft-values in Max-Ent RL.

188 3.2.1 Theory

189 Here we show that our objective is well behaved, considering the previously defined operator \mathbb{L}^β for
 190 random variables $\mathbb{L}^\beta(X) := \beta \log \mathbb{E} [e^{X/\beta}]$. First, we show it models the extremum.

191 **Lemma 3.1.** *For any $\beta_1 > \beta_2$, we have $\mathbb{L}^{\beta_1}(X) < \mathbb{L}^{\beta_2}(X)$. And $\mathbb{L}^\infty(X) = \mathbb{E}[X]$, $\mathbb{L}^0(X) =$
 192 $\sup(X)$. Thus, for any $\beta \in (0, \infty)$, the operator $\mathbb{L}^\beta(X)$ is a measure that interpolates between the
 193 expectation and the max of X .*

194 The operator $\mathbb{L}^\beta(X)$ is known as the cumulant-generating function or the log-Laplace transform, and
 195 is a measure of the tail-risk closely linked to the entropic value at risk (EVaR) [1].

196 **Lemma 3.2.** *The risk measure \mathcal{L} has a unique minima at $\beta \log \mathbb{E} [e^{X/\beta}]$. And an empirical risk $\hat{\mathcal{L}}$ is
 197 an unbiased estimate of the true risk. Furthermore, for $\beta \gg 1$, $\mathcal{L}(\theta) \approx \frac{1}{2\beta^2} \mathbb{E}_{x_i \sim \mathcal{D}} [(x_i - \theta)^2]$, thus
 198 behaving as the MSE loss with errors $\sim \mathcal{N}(0, \beta)$.*

199 In particular, the empirical loss $\hat{\mathcal{L}}$ over a dataset of N samples can be minimized using stochastic
 200 gradient-descent (SGD) methods to give an unbiased estimate of the LogSumExp over the N samples.

201 **Lemma 3.3.** *$\hat{\mathbb{L}}^\beta(X)$ over a finite N samples is a consistent estimator of the log-partition function
 202 $\mathbb{L}^\beta(X)$. Similarly, $\exp(\hat{\mathbb{L}}^\beta(X)/\beta)$ is an unbiased estimator for the partition function $Z = \mathbb{E} [e^{X/\beta}]$*

203 We provide concentration bounds for Lemma 3.3, and further theoretical discussion on Gumbel
 204 Regression in Appendix A.

⁴We add -1 to make the loss 0 for a perfect fit, as $e^x - x - 1 \geq 0$ with equality at $x = 0$.

3.3 MaxEnt RL without Entropy

Given Gumbel Regression can be used to directly model the LogSumExp, we apply it to Q-learning. First, we establish the connection of our framework with conservative Q-learning [19].

Lemma 3.4. *Consider the loss objective over Q-functions:*

$$\mathcal{L}(Q) = \mathbb{E}_{\mathbf{s} \sim \rho_\mu, \mathbf{a} \sim \mu(\cdot|\mathbf{s})} \left[e^{(\mathcal{T}^\pi \hat{Q}^k(\mathbf{s}, \mathbf{a}) - Q(\mathbf{s}, \mathbf{a}))/\beta} \right] - \mathbb{E}_{\mathbf{s} \sim \rho_\mu, \mathbf{a} \sim \pi(\cdot|\mathbf{s})} [(\mathcal{T}^\pi \hat{Q}^k(\mathbf{s}, \mathbf{a}) - Q(\mathbf{s}, \mathbf{a}))/\beta] - 1 \quad (8)$$

where $\mathcal{T}^\pi := r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}'|\mathbf{s}, \mathbf{a}} \mathbb{E}_{\mathbf{a}' \sim \pi} [Q(\mathbf{s}', \mathbf{a}')] is the vanilla Bellman operator under the policy $\pi(\mathbf{a}|\mathbf{s})$. Then minimizing \mathcal{L} gives the update rule:$

$$\forall \mathbf{s}, \mathbf{a}, k \quad \hat{Q}^{k+1}(\mathbf{s}, \mathbf{a}) = \mathcal{T}^\pi \hat{Q}^k(\mathbf{s}, \mathbf{a}) - \beta \log \frac{\pi(\mathbf{a}|\mathbf{s})}{\mu(\mathbf{a}|\mathbf{s})} = \mathcal{B}^\pi \hat{Q}^k(\mathbf{s}, \mathbf{a}).$$

Thus this transforms a regular Bellman backup into the soft-Bellman backup without the need of entropies.

Using the above lemma we can convert standard RL into Max-Ent RL. Here, $\mathcal{L}(\cdot)$ does a conservative Q-update similar to CQL [19] with the nice property that the implied conservative term is just the KL-constraint between π and μ .⁵ This enforces a entropy-regularization on our policy with respect to the behavior policy without the need of entropy. Thus, soft-Q learning naturally emerges as a conservative update on regular Q-learning under our objective. Here, Equation 8 is the dual of the KL-divergence between μ and π [12], and we motivate this objective for RL and establish formal equivalence with conservative Q-learning in Appendix C.

In our framework, we use the MaxEnt Bellman operator \mathcal{B}^* which gives our *ExtremeQ* loss:

$$\mathcal{L}(Q) = \mathbb{E}_{\mathbf{s}, \mathbf{a} \sim \mu} \left[e^{(\mathcal{B}^* \hat{Q}^k(\mathbf{s}, \mathbf{a}) - Q(\mathbf{s}, \mathbf{a}))/\beta} \right] - \mathbb{E}_{\mathbf{s}, \mathbf{a} \sim \mu} [(\mathcal{B}^* \hat{Q}^k(\mathbf{s}, \mathbf{a}) - Q(\mathbf{s}, \mathbf{a}))/\beta] - 1 \quad (9)$$

and is the same as our Gumbel loss from the previous section.

This gives an update rule: $\hat{Q}^{k+1}(\mathbf{s}, \mathbf{a}) = \mathcal{B}^* \hat{Q}^k(\mathbf{s}, \mathbf{a})$. $\mathcal{L}(\cdot)$ here requires estimation of \mathcal{B}^* which is very hard in continuous action spaces. Under deterministic dynamics, \mathcal{L} can be obtained without \mathcal{B}^* as shown in Appendix C. However, in general we still need to estimate \mathcal{B}^* . Next, we motivate how we can solve this issue.

Consider the soft-Bellman equation from Section 2.1 (Equation 1),

$$\mathcal{B}^* Q = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}' \sim P(\cdot|\mathbf{s}, \mathbf{a})} [V^*(\mathbf{s}')], \quad (10)$$

where $V^*(\mathbf{s}) = \mathbb{L}_{\mathbf{a} \sim \mu(\cdot|\mathbf{s})}^\beta [Q(\mathbf{s}, \mathbf{a})]$. Then V^* can be directly estimated using Gumbel regression by setting the temperature β to the regularization strength in the MaxEnt framework.

This gives us the following *ExtremeV* loss objective:

$$\mathcal{J}(V) = \mathbb{E}_{\mathbf{s}, \mathbf{a} \sim \mu} \left[e^{(\hat{Q}^k(\mathbf{s}, \mathbf{a}) - V(\mathbf{s}))/\beta} \right] - \mathbb{E}_{\mathbf{s}, \mathbf{a} \sim \mu} [(\hat{Q}^k(\mathbf{s}, \mathbf{a}) - V(\mathbf{s}))/\beta] - 1. \quad (11)$$

Lemma 3.5. *Minimizing \mathcal{J} over values gives the update rule: $\hat{V}^k(\mathbf{s}) = \mathbb{L}_{\mathbf{a} \sim \mu(\cdot|\mathbf{s})}^\beta [\hat{Q}^k(\mathbf{s}, \mathbf{a})]$.*

Then we can obtain V^* from $Q(\mathbf{s}, \mathbf{a})$ using Gumbel regression and substitute in Equation 10 to estimate the optimal bellman backup $\mathcal{B}^* Q$. Thus, Lemma 3.4 and 3.5 give us a scheme to solve the Max-Ent RL problem without the need of entropy.

3.4 Learning Policies

In the above section we derived a Q-learning strategy that does not require explicit use of a policy π . However, in continuous settings we still often want to recover a policy that can be run in the environment. Per Eq. 2 (Section 2.2), the optimal MaxEnt policy $\pi^*(\mathbf{a}|\mathbf{s}) = \mu(\mathbf{a}|\mathbf{s}) e^{(Q(\mathbf{s}, \mathbf{a}) - V(\mathbf{s}))/\beta}$. By minimizing the KL-divergence between π and the optimal π^* induced by Q and V we obtain the objective:

$$\pi^* = \underset{\pi}{\operatorname{argmin}} \mathbb{E}_{\rho_\mu(\mathbf{s}, \mathbf{a})} [e^{(Q(\mathbf{s}, \mathbf{a}) - V(\mathbf{s}))/\beta} \log \pi]. \quad (12)$$

⁵In fact, all theorems of CQL [19] hold for our objective by simple replacement of D_{CQL} with D_{KL} .

238 If we take ρ_μ to be a dataset \mathcal{D} generated from a behavior policy $\pi_{\mathcal{D}}$, we exactly recover the AWAC
 239 objective used by prior works in Offline RL [27]. This objective does not require sampling actions,
 240 which may potentially take $Q(s, a)$ out of distribution. Alternatively, if we want to sample from
 241 the policy we can minimize the Reverse-KL divergence which gives us the SAC-like actor update:

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}_{\rho_{\pi(s)\pi(a|s)}} [Q(s, \mathbf{a}) - \beta \log(\pi(\mathbf{a}|s)/\mu(\mathbf{a}|s))]. \quad (13)$$

242 Interestingly, we note this doesn't depend on $V(s)$. If μ is chosen to be the last policy π_k , the
 243 second term becomes the KL-divergence between the current policy and π_k , performing a trust
 244 region update on π [33, 37].⁶ While estimating the log ratio $\log(\pi(\mathbf{a}|s)/\mu(\mathbf{a}|s))$ can be difficult
 245 depending on choice of μ , we can remove the need for an analytic form of μ using our Gumbel loss
 246 objective. As applying our Gumbel Loss \mathcal{J} on samples from μ leads to soft-Q values of the form
 247 $Q(s, \mathbf{a}) - \beta \log(\pi(\mathbf{a}|s)/\mu(\mathbf{a}|s))$, our policy can be updated by simply maximizing the soft-Q values.
 248 Updating \mathcal{J} with samples from π_k thus naturally gives rise to trust-region style policy updates.

249 3.5 Practical Algorithms

250 In this section we develop a practical approach to
 251 Extreme Q-learning (\mathcal{X} -QL) for both online and of-
 252 fline RL. We consider parameterized functions $V_\theta(s)$,
 253 $Q_\phi(s, \mathbf{a})$, and $\pi_\psi(\mathbf{a}|s)$ and let \mathcal{D} be the training data
 254 distribution. A core issue with directly optimiz-
 255 ing Eq. 10 is over-optimism about dynamics [20]
 256 when using simple-sample estimates for the Bellman
 257 backup. To overcome this issue in stochastic settings,
 258 we separate out the optimization of V_θ from that of
 259 Q_ϕ following Section 3.3. We learn V_θ using Eq. 11
 260 to directly fit the optimal soft-values $V^*(s)$ based on
 261 Gumbel regression. Using $V_\theta(s')$ we can get single-
 262 sample estimates of B^* as $r(s, \mathbf{a}) + \gamma V_\theta(s')$. Now we can learn an unbiased expectation over the
 263 dynamics, $Q_\phi \approx \mathbb{E}_{s'|s, \mathbf{a}} [r(s, \mathbf{a}) + \gamma V_\theta(s')]$ by minimizing the Mean-squared-error (MSE) loss
 264 between the single-sample targets and Q_ϕ :

$$\mathcal{L}(\phi) = \mathbb{E}_{(s, \mathbf{a}, s') \sim \mathcal{D}} [(Q_\phi(s, \mathbf{a}) - r(s, \mathbf{a}) - \gamma V_\theta(s'))^2]. \quad (14)$$

265 In deterministic dynamics, our approach is largely simplified and we directly learn a single Q_ϕ using
 266 Eq. 9 without needing to learn B^* or V^* . Similarly, we learn soft-optimal policies using Eq. 12
 267 (offline) or Eq. 13 (online) settings.

268 **Offline RL.** In the offline setting, \mathcal{D} is specified as an offline dataset assumed to be collected with
 269 the behavior policy $\pi_{\mathcal{D}}$. Here, learning values with Eq. 11 has a number of practical benefits. First,
 270 we are able to fit the optimal soft-values V^* *without sampling from a policy network*, which has
 271 been shown to cause large out-of-distribution errors in the offline setting where mistakes cannot be
 272 corrected by collecting additional data. Second, we inherently enforce a KL-constraint on the optimal
 273 policy π^* and the behavior policy $\pi_{\mathcal{D}}$. This provides tunable conservatism via the temperature β .
 274 After offline training of Q_ϕ and V_θ , we can recover the policy post-training using the AWAC objective
 275 (Eq. 12). Our practical implementation follows the training style of Kostrikov et al. [17], but we train
 276 value network using using our ExtremeQ loss.

277 **Online RL.** In the online setting, \mathcal{D} is usually given as a replay buffer of previously sampled states
 278 and actions. In practice, however, obtaining a good estimate of $V^*(s')$ requires that we sample
 279 actions with high Q-values instead of uniform sampling from \mathcal{U} . As online learning allows agents to
 280 correct over-optimism by collecting additional data, we use a previous version of the policy network
 281 π_{ψ} to sample actions for the Bellman backup, amounting to the trust-region policy update at the end
 282 of Section 3.4. In practice, we modify SAC and TD3 with our formulation. To embue SAC [14] with
 283 the benefits of Extreme Q-learning, we simply train V_θ using Eq. 11 with $s \sim \mathcal{D}$, $\mathbf{a} \sim \pi_{\psi_k}(\mathbf{a}|s)$. TD3
 284 by default doesn't use a value network, and thus we use our algorithm for deterministic dynamics by
 285 changing the loss to train Q in TD3 to directly follow Eq. 9.

⁶Choosing μ to be uniform \mathcal{U} gives the regular SAC update.

4 Experiments

We compare our Extreme Q-Learning (\mathcal{X} -QL) approach to existing works obtaining state-of-the-art results across a wide set of continuous control tasks in both online and offline settings. More details, offline results on Adroit, and hyper-parameter settings can be found in Appendix E.

4.1 Offline RL

Table 1: Averaged normalized scores on MuJoCo locomotion and Ant Maze tasks.

	Dataset	BC	10%BC	DT	AWAC	Onestep RL	TD3+BC	CQL	IQL	\mathcal{X} -QL	(Gain)
Gym	halfcheetah-medium-v2	42.6	42.5	42.6	43.5	48.4	48.3	44.0	47.4	48.3	-0.1
	hopper-medium-v2	52.9	56.9	67.6	57.0	59.6	59.3	58.5	66.3	75.6	+8.0
	walker2d-medium-v2	75.3	75.0	74.0	72.4	81.8	83.7	72.5	78.3	86.0	+4.2
	halfcheetah-medium-replay-v2	36.6	40.6	36.6	40.5	38.1	44.6	45.5	44.2	45.4	-0.1
	hopper-medium-replay-v2	18.1	75.9	82.7	37.2	97.5	60.9	95.0	94.7	102.4	+4.9
	walker2d-medium-replay-v2	26.0	62.5	66.6	27.0	49.5	81.8	77.2	73.9	86.1	+4.3
	halfcheetah-medium-expert-v2	55.2	92.9	86.8	42.8	93.4	90.7	91.6	86.7	94.2	+0.8
	hopper-medium-expert-v2	52.5	110.9	107.6	55.8	103.3	98.0	105.4	91.5	112.3	+1.4
AntMaze	walker2d-medium-expert-v2	107.5	109.0	108.1	74.5	113.0	110.1	108.8	109.6	112.7	-0.3
	antmaze-umaze-v0	54.6	62.8	59.2	56.7	64.3	78.6	74.0	87.5	97.1	+9.6
	antmaze-umaze-diverse-v0	45.6	50.2	53.0	49.3	60.7	71.4	84.0	62.2	85.2	+1.2
	antmaze-medium-play-v0	0.0	5.4	0.0	0.0	0.3	10.6	61.2	71.2	83.2	+12.0
	antmaze-medium-diverse-v0	0.0	9.8	0.0	0.7	0.0	3.0	53.7	70.0	78.5	+8.5
	antmaze-large-play-v0	0.0	0.0	0.0	0.0	0.0	0.2	15.8	39.6	48.7	+9.1
Franka	antmaze-large-diverse-v0	0.0	6.0	0.0	1.0	0.0	0.0	14.9	47.5	53.1	+5.6
	kitchen-complete-v0	65.0	-	-	-	-	-	43.8	62.5	82.4	+17.4
	kitchen-partial-v0	38.0	-	-	-	-	-	49.8	46.3	75.0	+28.7
runtime	kitchen-mixed-v0	51.5	-	-	-	-	-	51.0	51.0	62.5	+11.0
		10m	10m	960m	20m	20m	20m	80m	20m	10m*	

*We see very fast convergence for our method and training for 0.5M steps suffices on most tasks (others are with 1M steps).

Our offline approach strongly outperforms prior methods [6, 18, 19, 17, 9] by 10-20 absolute points in many environments reaching the new *state-of-the-art* on the D4RL benchmark, as shown in Table 1. We particularly see large improvements on the AntMaze tasks, which require a significant amount of “stitching” between trajectories [17] and get **double-digit** improvement on the challenging Franka tasks. We find performance on the Gym locomotion tasks to be already largely saturated. Moreover, our method convergences significantly faster than IQL, and we show learning curves in Appendix D. \mathcal{X} -QL can be easily fine-tuned using online data to attain even higher performance as shown in Table 2, surpassing the final performance of prior works by a wide margin. On a variety of AntMaze tasks, our offline results *before* fine-tuning are on par-with IQL’s performance *after* online fine-tuning. Like Kostrikov et al. [17] we tune β for different environments, which corresponds to the weight on our conservative KL penalty.

4.2 Online RL

We compare ExtremeQ variants of SAC [14] and TD3 [10], denoted \mathcal{X} -SAC and \mathcal{X} -TD3, to their vanilla versions on tasks in the DM Control, shown in Figure 3. Across all tasks an ExtremeQ variant matches or surpasses the performance of baselines.

We see particularly large gains in the Hopper environment, and more significant gains in comparison to TD3 overall. Consistent with SAC [14], we find the temperature β needs to be tuned for different environments with different reward scales and sparsity. A core component of TD3 introduced by [10] is Double Q-Learning, which takes the minimum of two Q functions to remove overestimate bias in the Q -target. As we assume errors to be Gumbel distributed, we expect our \mathcal{X} -variants to be more robust to such errors. In all environments except Cheetah Run, our \mathcal{X} -TD3 without the Double-Q trick performs better than standard TD3. Interestingly, we find that in the DM Control environments Double-Q learning did not boost performance across the board. While the gains from applying Extreme-Q learning are modest in the online setting, none of our methods require access to the policy distribution $\pi(a|s)$. In particular, \mathcal{X} -SAC only requires samples from π_ψ , unlike regular SAC which incorporates the log probabilities of π_ψ into value estimates.

Table 2: Finetuning results on the AntMaze environments

Dataset	CQL	IQL	\mathcal{X} -QL
umaze-v0	70.1 → 99.4	86.7 → 96.0	96.1 → 99.6
umaze-diverse-v0	31.1 → 99.4	75.0 → 84.0	82.5 → 99.0
medium-play-v0	23.0 → 0.0	72.0 → 95.0	80.2 → 97.0
medium-diverse-v0	23.0 → 32.3	68.3 → 92.0	74.6 → 97.1
large-play-v0	1.0 → 0.0	25.5 → 46.0	45.1 → 59.3
large-diverse-v0	1.0 → 0.0	42.6 → 60.7	52.2 → 82.1

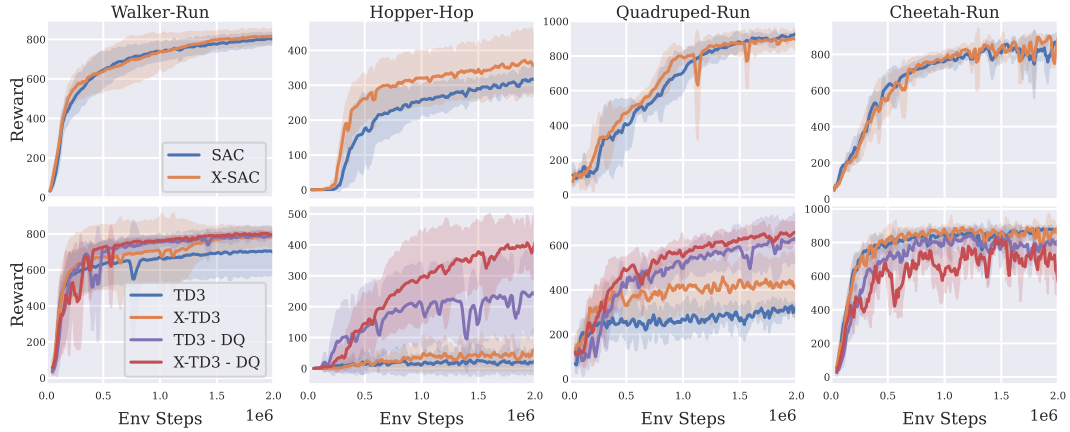


Figure 3: Results on the DM Control for SAC and TD3 based versions of ExtremeQ Learning.

5 Related Work

Our approach builds on top of literature in online & offline RL, and here we review the most salient works. It is worth noting that the inspiration for our novel framework comes from works on choice theory in econometrics [32, 24], and our Gumbel loss was motivated by the work of IQ-Learn [12].

Online RL. Our work bridges the theoretical gap between RL and Max-Ent RL by introducing our Gumbel loss function. Unlike past work in MaxEnt RL [14, 7], our method does not require explicit entropy estimation and instead addresses the problem of obtaining soft-value estimates (LogSumExp) in high-dimensional or continuous spaces [38] by directly modeling them via our proposed Gumbel loss, which to our knowledge has not previously been used in RL. Nevertheless, our loss objective is intrinsically linked to the KL divergence, and similar objectives have been used for mutual information estimation [31]. IQ-Learn [12] - an avant-garde approach to Imitation Learning (IL) - introduced the same loss in IL to obtain an unbiased dual form for the reverse KL-divergence between an expert and policy distribution. Prior work in RL has also examined using other types of loss functions [3] or other formulations of the argmax in order to ease optimization [2]. Distinct from most off-Policy RL Methods [22, 10, 14], we directly model \mathcal{B}^* like [13, 16] but attain significantly more stable results.

Offline RL. Prior works in offline RL can largely be categorized as relying on constrained or regularized Q-learning [39, 9, 11, 18, 19, 27], or extracting a greedy policy from the known behavior policy [30, 5, 6]. Most similar to our work, IQL [17] fits expectiles of the Q-function of the behavior policy, but is not motivated to solve a particular problem or remain conservative. On the other hand, conservatism in CQL [19] is motivated by lower-bounding the Q-function’s value. Our method shares the best of both worlds – like IQL we do not evaluate the Q-function on potentially out of distribution actions and like CQL we enjoy the benefits of conservatism. Compared to CQL, our approach uses a KL constraint with the behavior policy, and for the first time successfully extends soft-Q learning to offline RL without needing a policy or explicit entropy values.

6 Conclusion

We propose Extreme Q-Learning, a new framework for MaxEnt RL that directly estimates the optimal Bellman backup \mathcal{B}^* without relying on explicit access to a policy. Theoretically, we bridge the gap between the regular, soft, and conservative Q-learning formulations. Empirically, we show that our framework can be used to develop simple SOTA RL algorithms. A number of future directions remain – improving stability in training with exponential loss functions, better selection of policy distributions μ in the online setting, and integrating automatic tuning methods for β . Finally, we hope that our framework can find uses beyond RL in general Machine Learning using our Gumbel loss for estimation of log-partition functions.

References

- [1] A. Ahmadi-Javid. Entropic value-at-risk: A new coherent risk measure. *Journal of Optimization Theory and Applications*, 155(3):1105–1123, 2012. URL <https://EconPapers.repec.org/>

RePEc:spr:joptap:v:155:y:2012:i:3:d:10.1007_s10957-011-9968-2. 5

[2] Kavosh Asadi and Michael L Littman. An alternative softmax operator for reinforcement learning. In *International Conference on Machine Learning*, pages 243–252. PMLR, 2017. 9

[3] Joan Bas-Serrano, Sebastian Curi, Andreas Krause, and Gergely Neu. Logistic q-learning. In *International Conference on Artificial Intelligence and Statistics*, pages 3610–3618. PMLR, 2021. 9

[4] M. Bloem and N. Bambos. Infinite time horizon maximum causal entropy inverse reinforcement learning. *53rd IEEE Conference on Decision and Control*, pages 4911–4916, 2014. 2

[5] David Brandfonbrener, Will Whitney, Rajesh Ranganath, and Joan Bruna. Offline rl without off-policy evaluation. *Advances in Neural Information Processing Systems*, 34:4933–4946, 2021. 9

[6] Lili Chen, Kevin Lu, Aravind Rajeswaran, Kimin Lee, Aditya Grover, Misha Laskin, Pieter Abbeel, Aravind Srinivas, and Igor Mordatch. Decision transformer: Reinforcement learning via sequence modeling. *Advances in neural information processing systems*, 34, 2021. 8, 9

[7] Benjamin Eysenbach and Sergey Levine. If maxent {rl} is the answer, what is the question?, 2020. URL <https://openreview.net/forum?id=SkxcZCNKDS>. 9

[8] R. A. Fisher and L. H. C. Tippett. Limiting forms of the frequency distribution of the largest or smallest member of a sample. *Mathematical Proceedings of the Cambridge Philosophical Society*, 24(2):180–190, 1928. doi: 10.1017/S0305004100015681. 3

[9] Scott Fujimoto and Shixiang Shane Gu. A minimalist approach to offline reinforcement learning. *Advances in Neural Information Processing Systems*, 34, 2021. 8, 9

[10] Scott Fujimoto, Herke van Hoof, and David Meger. Addressing function approximation error in actor-critic methods. *ArXiv*, abs/1802.09477, 2018. 1, 8, 9, 15, 21

[11] Scott Fujimoto, David Meger, and Doina Precup. Off-policy deep reinforcement learning without exploration. In *International Conference on Machine Learning*, pages 2052–2062. PMLR, 2019. 9

[12] Divyansh Garg, Shuvam Chakraborty, Chris Cundy, Jiaming Song, and Stefano Ermon. Iq-learn: Inverse soft-q learning for imitation. In *Thirty-Fifth Conference on Neural Information Processing Systems*, 2021. URL <https://openreview.net/forum?id=Aeo-xqtb5p>. 6, 9, 20

[13] Tuomas Haarnoja, Haoran Tang, Pieter Abbeel, and Sergey Levine. Reinforcement learning with deep energy-based policies. 2017. 3, 9, 21

[14] Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. In *International conference on machine learning*, pages 1861–1870. PMLR, 2018. 1, 3, 4, 7, 8, 9

[15] Tamir Hazan and Tommi Jaakkola. On the partition function and random maximum a-posteriori perturbations. *arXiv preprint arXiv:1206.6410*, 2012. 3

[16] Nicolas Heess, Gregory Wayne, David Silver, Timothy Lillicrap, Tom Erez, and Yuval Tassa. Learning continuous control policies by stochastic value gradients. *Advances in neural information processing systems*, 28, 2015. 9

[17] Ilya Kostrikov, Ashvin Nair, and Sergey Levine. Offline reinforcement learning with implicit q-learning. *arXiv preprint arXiv:2110.06169*, 2021. 7, 8, 9, 12, 21

[18] Aviral Kumar, Justin Fu, Matthew Soh, George Tucker, and Sergey Levine. Stabilizing off-policy q-learning via bootstrapping error reduction. *Advances in Neural Information Processing Systems*, 32, 2019. 8, 9

- [19] Aviral Kumar, Aurick Zhou, George Tucker, and Sergey Levine. Conservative q-learning for offline reinforcement learning. *Advances in Neural Information Processing Systems*, 33: 1179–1191, 2020. 1, 6, 8, 9, 19
- [20] Sergey Levine. Reinforcement learning and control as probabilistic inference: Tutorial and review. *arXiv preprint arXiv:1805.00909*, 2018. 7
- [21] Qing Li. Continuous control benchmark of deepmind control suite and mujoco. https://github.com/LQNew/Continuous_Control_Benchmark, 2021. 22
- [22] Timothy P Lillicrap, Jonathan J Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. *arXiv preprint arXiv:1509.02971*, 2015. 9
- [23] R.Duncan Luce. The choice axiom after twenty years. *Journal of Mathematical Psychology*, 15 (3):215–233, 1977. ISSN 0022-2496. doi: [https://doi.org/10.1016/0022-2496\(77\)90032-3](https://doi.org/10.1016/0022-2496(77)90032-3). URL <https://www.sciencedirect.com/science/article/pii/0022249677900323>. 3, 15
- [24] Daniel McFadden. Conditional logit analysis of qualitative choice behavior. 1972. 2, 3, 9, 14, 15
- [25] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller. Playing atari with deep reinforcement learning. *arXiv preprint arXiv:1312.5602*, 2013. 1
- [26] Alexander McFarlane Mood. Introduction to the theory of statistics. 1950. 3
- [27] Ashvin Nair, Abhishek Gupta, Murtaza Dalal, and Sergey Levine. Awac: Accelerating online reinforcement learning with offline datasets. *arXiv preprint arXiv:2006.09359*, 2020. 7, 9
- [28] Gergely Neu, Anders Jonsson, and V. Gómez. A unified view of entropy-regularized markov decision processes. *ArXiv*, abs/1705.07798, 2017. 3
- [29] George Papandreou and Alan L Yuille. Gaussian sampling by local perturbations. *Advances in Neural Information Processing Systems*, 23, 2010. 3
- [30] Xue Bin Peng, Aviral Kumar, Grace Zhang, and Sergey Levine. Advantage-weighted regression: Simple and scalable off-policy reinforcement learning. *arXiv preprint arXiv:1910.00177*, 2019. 9
- [31] Ben Poole, Sherjil Ozair, Aäron van den Oord, Alexander A. Alemi, and G. Tucker. On variational bounds of mutual information. In *ICML*, 2019. 9
- [32] John Rust. Structural estimation of markov decision processes. In R. F. Engle and D. McFadden, editors, *Handbook of Econometrics*, volume 4, chapter 51, pages 3081–3143. Elsevier, 1 edition, 1986. URL https://editorialexpress.com/jrust/papers/handbook_ec_v4_rust.pdf. 2, 3, 9
- [33] John Schulman, Sergey Levine, Pieter Abbeel, Michael Jordan, and Philipp Moritz. Trust region policy optimization. In *International conference on machine learning*, pages 1889–1897. PMLR, 2015. 1, 7
- [34] Slavko Simić. On a new converse of jensen’s inequality. *Publications De L’institut Mathématique*, 85:107–110, 01 2009. doi: 10.2298/PIM0999107S. 17
- [35] Yuval Tassa, Yotam Doron, Alistair Muldal, Tom Erez, Yazhe Li, Diego de Las Casas, David Budden, Abbas Abdolmaleki, Josh Merel, Andrew LeFrancq, et al. Deepmind control suite. *arXiv preprint arXiv:1801.00690*, 2018. 4
- [36] Sebastian Thrun and Anton Schwartz. Issues in using function approximation for reinforcement learning. 1999. 15
- [37] Nino Vieillard, Tadashi Kozuno, Bruno Scherrer, Olivier Pietquin, Rémi Munos, and Matthieu Geist. Leverage the average: an analysis of kl regularization in rl. *34th Conference on Neural Information Processing Systems*, 2020. 7

- [38] Nino Vieillard, Marcin Andrychowicz, Anton Raichuk, Olivier Pietquin, and Matthieu Geist. Implicitly regularized rl with implicit q-values. *arXiv preprint arXiv:2108.07041*, 2021. 9
- [39] Yifan Wu, George Tucker, and Ofir Nachum. Behavior regularized offline reinforcement learning. *arXiv preprint arXiv:1911.11361*, 2019. 9
- [40] Denis Yarats and Ilya Kostrikov. Soft actor-critic (sac) implementation in pytorch. https://github.com/denisyarats/pytorch_sac, 2020. 21, 22
- [41] G. Alastair Young. High-dimensional statistics: A non-asymptotic viewpoint, martin j. wainwright, cambridge university press, 2019, xvii 552 pages, £57.99, hardback isbn: 978-1-1084-9802-9. *International Statistical Review*, 88(1):258–261, 2020. doi: <https://doi.org/10.1111/insr.12370>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1111/insr.12370>. 3
- [42] Brian D Ziebart. *Modeling purposeful adaptive behavior with the principle of maximum causal entropy*. Carnegie Mellon University, 2010. 1

Checklist

The checklist follows the references. Please read the checklist guidelines carefully for information on how to answer these questions. For each question, change the default **[TODO]** to **[Yes]**, **[No]**, or **[N/A]**. You are strongly encouraged to include a **justification to your answer**, either by referencing the appropriate section of your paper or providing a brief inline description. For example:

- Did you include the license to the code and datasets? **[Yes]** See Section ??.
- Did you include the license to the code and datasets? **[No]** The code and the data are proprietary.
- Did you include the license to the code and datasets? **[N/A]**

Please do not modify the questions and only use the provided macros for your limit. In your paper, please delete this instructions block and only keep the Checklist section heading above along with the questions/answers below.

1. For all authors...

- (a) Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope? **[Yes]** See the results section and intro/abstract
- (b) Did you describe the limitations of your work? **[Yes]** See the conclusion, section 6.
- (c) Did you discuss any potential negative societal impacts of your work? **[Yes]** See the appendix
- (d) Have you read the ethics review guidelines and ensured that your paper conforms to them? **[Yes]**

2. If you are including theoretical results...

- (a) Did you state the full set of assumptions of all theoretical results? **[Yes]** See sections 2 and 3
- (b) Did you include complete proofs of all theoretical results? **[Yes]** We have largely expanded on results included in Section 2 and 3 in the Appendix.

3. If you ran experiments...

- (a) Did you include the code, data, and instructions needed to reproduce the main experimental results (either in the supplemental material or as a URL)? **[Yes]** Supplementary material
- (b) Did you specify all the training details (e.g., data splits, hyperparameters, how they were chosen)? **[Yes]** See the last section of the appendix
- (c) Did you report error bars (e.g., with respect to the random seed after running experiments multiple times)? **[Yes]** Error bars are shown for Online RL, for offline RL we report the average over seeds as done in [17] which we compare to.

- 496 (d) Did you include the total amount of compute and the type of resources used (e.g., type
497 of GPUs, internal cluster, or cloud provider)? [Yes] See the last section of the appendix.
- 498 4. If you are using existing assets (e.g., code, data, models) or curating/releasing new assets...
- 499 (a) If your work uses existing assets, did you cite the creators? [Yes] We cited DM Control,
500 Open AI Gym, D4RL, etc.
- 501 (b) Did you mention the license of the assets? [N/A] All assets have open licences
- 502 (c) Did you include any new assets either in the supplemental material or as a URL? [No]
- 503 (d) Did you discuss whether and how consent was obtained from people whose data you're
504 using/curating? [N/A]
- 505 (e) Did you discuss whether the data you are using/curating contains personally identifiable
506 information or offensive content? [N/A]
- 507 5. If you used crowdsourcing or conducted research with human subjects...
- 508 (a) Did you include the full text of instructions given to participants and screenshots, if
509 applicable? [N/A]
- 510 (b) Did you describe any potential participant risks, with links to Institutional Review
511 Board (IRB) approvals, if applicable? [N/A]
- 512 (c) Did you include the estimated hourly wage paid to participants and the total amount
513 spent on participant compensation? [N/A]

514 A Extreme Theory of Q-learning

515 In this section, we functionally analyze Q-learning using our framework and further develop the
516 Gumbel Error Model (GEM) for MDPs.

517 A.1 Rust-McFadden Model of MDPs

518 For an MDP following the Bellman equations, we assume the observed rewards to be stochastic due
519 to an unobserved component of the state. Let \mathbf{s} be the observed state, and (\mathbf{s}, \mathbf{z}) be the actual state
520 with hidden component \mathbf{z} . Then,

$$Q(\mathbf{s}, \mathbf{z}, \mathbf{a}) = R(\mathbf{s}, \mathbf{z}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}' \sim P(\cdot | \mathbf{s}, \mathbf{a})} [\mathbb{E}_{\mathbf{z}' | \mathbf{s}'} [V(\mathbf{s}', \mathbf{z}')]], \quad (15)$$

$$V(\mathbf{s}, \mathbf{z}) = \max_{\mathbf{a}} Q(\mathbf{s}, \mathbf{z}, \mathbf{a}). \quad (16)$$

521 **Lemma A.1.** *Given, 1) conditional independence (CI) assumption that \mathbf{z}' depends only on \mathbf{s}' , i.e.*
522 *$p(\mathbf{s}', \mathbf{z}' | \mathbf{s}, \mathbf{z}, \mathbf{a}) = p(\mathbf{z}' | \mathbf{s}') p(\mathbf{s}' | \mathbf{s}, \mathbf{a})$ and 2) additive separability (AS) assumption on the hidden noise:*
523 *$R(\mathbf{s}, \mathbf{a}, \mathbf{z}) = r(\mathbf{s}, \mathbf{a}) + \epsilon(\mathbf{z}, \mathbf{a})$.*

524 *Then for i.i.d. $\epsilon(\mathbf{z}, \mathbf{a}) \sim \mathcal{G}(0, \beta)$, we recover the soft-Bellman equations for $Q(\mathbf{s}, \mathbf{z}, \mathbf{a}) = q(\mathbf{s}, \mathbf{a}) +$*
525 *$\epsilon(\mathbf{z}, \mathbf{a})$ and $v(\mathbf{s}) = \mathbb{E}_{\mathbf{z}} [V(\mathbf{s}, \mathbf{z})]$, with rewards $r(\mathbf{s}, \mathbf{a})$ and entropy regularization β .*

526 *Hence, a soft-MDP in MaxEntRL is equivalent to an MDP with an extra hidden variable in the state*
527 *that introduces i.i.d. Gumbel noise in the rewards and follows the AS+CI conditions.*

528 *Proof.* We have,

$$q(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}' \sim P(\cdot | \mathbf{s}, \mathbf{a})} [\mathbb{E}_{\mathbf{z}' | \mathbf{s}'} [V(\mathbf{s}', \mathbf{z}')]] \quad (17)$$

$$v(\mathbf{s}) = \mathbb{E}_{\mathbf{z}} [V(\mathbf{s}, \mathbf{z})] = \mathbb{E}_{\mathbf{z}} [\max_{\mathbf{a}} (q(\mathbf{s}, \mathbf{a}) + \epsilon(\mathbf{z}, \mathbf{a}))]. \quad (18)$$

529 From this, we can get fixed-point equations for q and π ,

$$q(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}' \sim P(\cdot | \mathbf{s}, \mathbf{a})} [\mathbb{E}_{\mathbf{z}' | \mathbf{s}'} [\max_{\mathbf{a}'} (q(\mathbf{s}', \mathbf{a}') + \epsilon(\mathbf{z}', \mathbf{a}'))]], \quad (19)$$

$$\pi(\cdot | \mathbf{s}) = \mathbb{E}_{\mathbf{z}} [\arg\max_{\mathbf{a}} (q(\mathbf{s}, \mathbf{a}) + \epsilon(\mathbf{z}, \mathbf{a}))] \in \Delta_{\mathcal{A}}, \quad (20)$$

530 where $\Delta_{\mathcal{A}}$ is the set of all policies.

531 Now, let $\epsilon(\mathbf{z}, \mathbf{a}) \sim \mathcal{G}(0, \beta)$ and assumed independent for each (\mathbf{z}, \mathbf{a}) (or equivalently (\mathbf{s}, \mathbf{a}) due to
532 the CI condition). Then we can use the Gumbel-Max trick to recover the soft-Bellman equations for
533 $q(\mathbf{s}, \mathbf{a})$ and $v(\mathbf{s})$ with rewards $r(\mathbf{s}, \mathbf{a})$:

$$q(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}' \sim P(\cdot | \mathbf{s}, \mathbf{a})} [\mathbb{L}_{\mathbf{a}'}^{\beta} [q(\mathbf{s}', \mathbf{a}'))]], \quad (21)$$

$$\pi(\cdot | \mathbf{s}) = \text{softmax}_{\mathbf{a}} (q(\mathbf{s}, \mathbf{a})). \quad (22)$$

534 Thus, we have that the soft-Bellman optimality equation and related optimal policy can arise either
535 from the entropic regularization viewpoint or from the Gumbel error viewpoint for an MDP.

536 **Corollary A.1.1.** *Converse: An MDP following the Bellman optimality equation and having a*
537 *policy that is softmax distributed, necessarily has any i.i.d. noise in the rewards due to hidden state*
538 *variables be Gumbel distributed, given the AS+CI conditions hold.*

539 *Proof.* McFadden [24] proved this converse in his seminal work on discrete choice theory, that for
540 i.i.d. ϵ satisfying Equation 19 with a choice policy $\pi \sim \text{softmax}$ has ϵ be Gumbel distributed. And
541 we show a proof here similar to the original for MDPs.

542 Considering Equation 20, we want $\pi(\mathbf{a} | \mathbf{s})$ to be softmax distributed. Let ϵ have an unknown CDF F
543 and we consider there to be N possible actions. Then,

$$\begin{aligned} P(\arg\max_{\mathbf{a}} (q(\mathbf{s}, \mathbf{a}) + \epsilon(\mathbf{z}, \mathbf{a})) = \mathbf{a}_i | \mathbf{s}, \mathbf{z}) &= P(q(\mathbf{s}, \mathbf{a}_i) + \epsilon(\mathbf{z}, \mathbf{a}_i) \geq q(\mathbf{s}, \mathbf{a}_j) + \epsilon(\mathbf{z}, \mathbf{a}_j) \forall i \neq j | \mathbf{s}, \mathbf{z}) \\ &= P(\epsilon(\mathbf{z}, \mathbf{a}_j) - \epsilon(\mathbf{z}, \mathbf{a}_i) \leq q(\mathbf{s}, \mathbf{a}_i) - q(\mathbf{s}, \mathbf{a}_j) \forall i \neq j | \mathbf{s}, \mathbf{z}) \end{aligned}$$

544 Simplifying the notation, we write $\epsilon(\mathbf{z}, \mathbf{a}_i) = \epsilon_i$ and $q(\mathbf{s}, \mathbf{a}_i) = q_i$. Then $\epsilon_1, \dots, \epsilon_N$ has a joint CDF
 545 G :

$$G(\epsilon_1, \dots, \epsilon_N) = \prod_{j=1}^N P(\epsilon_j \leq \epsilon_i + q_i - q_j) = \prod_{j=1}^N F(\epsilon_i + q_i - q_j)$$

546 and we can get the required probability $\pi(i)$ as:

$$\pi(i) = \int_{\epsilon=-\infty}^{+\infty} \prod_{j=1, j \neq i}^N F(\epsilon + q_i - q_j) dF(\epsilon) \quad (23)$$

547 For $\pi = \text{softmax}(q)$, McFadden [24] proved the uniqueness of F to be the Gumbel CDF, assuming
 548 translation completeness property to hold for F . Later this uniqueness was shown to hold in general
 549 for any $N \geq 3$ [23]. \square

550 A.2 Gumbel Error Model (GEM) for MDPs

551 To develop our Gumbel Error Model (GEM) for MDPs under functional approximation as in Sec-
 552 tion 3.1, we follow our simplified scheme of M independent estimators \hat{Q} , which results in the
 553 following equation over $\bar{Q} = \mathbb{E}[\hat{Q}]$:

$$\bar{Q}_{t+1}(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}'|\mathbf{s}, \mathbf{a}} [\mathbb{E}_{\epsilon_t} [\max_{\mathbf{a}'} (\bar{Q}_t(\mathbf{s}', \mathbf{a}') + \epsilon_t(\mathbf{s}', \mathbf{a}'))]]. \quad (24)$$

554 Here, the maximum of random variables will generally be greater than the true max, i.e.
 555 $\mathbb{E}_{\epsilon} [\max_{\mathbf{a}'} (\bar{Q}(\mathbf{s}', \mathbf{a}') + \epsilon(\mathbf{s}', \mathbf{a}'))] \geq \max_{\mathbf{a}'} \bar{Q}(\mathbf{s}', \mathbf{a}')$ [36]. As a result, even initially zero-mean error
 556 can cause Q updates to propagate consistent overestimation bias through the Bellman equation. This
 557 is a known issue with function approximation in RL [10].

558 Now, we can use the Rust-McFadden model from before. To account for the stochasticity, we consider
 559 extra unobserved state variables z in the MDP to be the model parameters θ used in the functional
 560 approximation. The errors from functional approximation ϵ_t can thus be considered as noise added in
 561 the reward. Here, CI condition holds as ϵ is separate from the dynamics and becomes conditionally
 562 independent for each state-action pair and AS condition is implied. Then for \bar{Q} satisfying Equation 24,
 563 we can apply the McFadden-Rust model, which implies that for the policy to be soft-optimal i.e. a
 564 softmax over \bar{Q} , ϵ will be Gumbel distributed.

565 Conversely, for the i.i.d. $\epsilon \sim \mathcal{G}$, $\bar{Q}(\mathbf{s}, \mathbf{a})$ follows the soft-Bellman equations and $\pi(\mathbf{a}|\mathbf{s}) =$
 566 $\text{softmax}(\bar{Q}(\mathbf{s}, \mathbf{a}))$.

567 This indicates an optimality condition on the MDP – for us to eventually attain the optimal softmax
 568 policy in the presence of functional bootstrapping (Equation 24), the errors should follow the Gumbel
 569 distribution.

570 A.2.1 Time Evolution of Errors in MDPs under Deterministic dynamics

571 In this section, we characterize the time evolution of errors in an MDP using GEM. We assume
 572 deterministic dynamics to simplify our analysis.

573 We suppose that we know the distribution of Q -values at time t and model the evolution of this
 574 distribution through the Bellman equations. Let $Z_t(\mathbf{s}, \mathbf{a})$ be a random variable sampled from the
 575 distribution of Q -values at time t , then the following Bellman equation holds:

$$Z_{t+1}(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \max_{\mathbf{a}'} Z_t(\mathbf{s}', \mathbf{a}'). \quad (25)$$

576 Here, $Z_{t+1}(\mathbf{s}, \mathbf{a}) = \max_{\mathbf{a}'} [r(\mathbf{s}, \mathbf{a}) + \gamma Z_t(\mathbf{s}', \mathbf{a}')] is a maximal distribution and based on EVT should
 577 eventually converge to an extreme value distribution, which we can model as a Gumbel.$

578 Concretely, let's assume that we fix $Z_t(\mathbf{s}, \mathbf{a}) \sim \mathcal{G}(Q_t(\mathbf{s}, \mathbf{a}), \beta)$ for some $Q_t(\mathbf{s}, \mathbf{a}) \in \mathbb{R}$ and $\beta > 0$.
 579 Furthermore, we assume that the Q -value distribution is jointly independent over different state-
 580 actions i.e. $Z(\mathbf{s}, \mathbf{a})$ is independent from $Z(\mathbf{s}', \mathbf{a}')$ for $\forall (\mathbf{s}, \mathbf{a}) \neq (\mathbf{s}', \mathbf{a}')$. Then $\max_{\mathbf{a}'} Z_t(\mathbf{s}', \mathbf{a}') \sim$
 581 $\mathcal{G}(V(\mathbf{s}'), \beta)$ with $V(\mathbf{s}) = \mathbb{L}_{\mathbf{a}}^{\beta}[Q(\mathbf{s}, \mathbf{a})]$ using the Gumbel-max trick.

Then substituting in Equation 25 and rescaling Z_t with γ , we get:

$$Z_{t+1}(\mathbf{s}, \mathbf{a}) \sim \mathcal{G} \left(r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{L}_{\mathbf{a}'}^\beta [Q(\mathbf{s}', \mathbf{a}')], \gamma\beta \right). \quad (26)$$

So very interestingly the Q-distribution becomes a Gumbel process, where the location parameter $Q(\mathbf{s}, \mathbf{a})$ follows the optimal soft-Bellman equation. Similarly, the temperature scales as $\gamma\beta$ and the distribution becomes sharper after every timestep.

After a number of timesteps, we see that $Z(\mathbf{s}, \mathbf{a})$ eventually collapses to the Delta distribution over the unique contraction $Q^*(\mathbf{s}, \mathbf{a})$. Here, γ controls the rate of decay of the Gumbel distribution into the collapsed Delta distribution. Thus we get the expected result in deterministic dynamics that the optimal Q-function will be deterministic and its distribution will be peaked.

So if a Gumbel error enters into the MDP through a functional error or some other source at a timestep t in some state s , it will trigger off a wave that propagates the Gumbel error into its child states following Equation 26. Thus, this Gumbel error process will decay at a γ rate every timestep and eventually settle down with Q-values reaching the the steady solution Q^* . The variance of this Gumbel process given as $\frac{\pi^2}{6}\beta^2$ will decay as γ^2 , similarly the bias will decay as γ -contraction in the \mathcal{L}^∞ norm.

Hence, GEM gives us an analytic characterization of error propogation in MDPs under deterministic dynamics.

Neertheless under stochastic dynamics, characterization of errors using GEM becomes non-trivial as Gumbel is not mean-stable unlike the Gaussian distribution. We hypothesise that the errors will follow some mix of Gumbel-Gaussian distributions, and leave this characterization as a future open direction.

A.3 Extreme Q-Learning (\mathcal{X} -QL)

For the soft-Bellman equation given as:

$$Q(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}' \sim P(\cdot|\mathbf{s}, \mathbf{a})} V(\mathbf{s}), \quad (27)$$

$$V(\mathbf{s}) = \mathbb{L}_{\mu(\cdot|\mathbf{s})}^\beta (Q(\mathbf{s}, \mathbf{a})), \quad (28)$$

we have the fixed-point characterization, that can be found with a recurrence:

$$V(\mathbf{s}) = \mathbb{L}_{\mu(\cdot|\mathbf{s})}^\beta \left(r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{E}_{\mathbf{s}' \sim P(\cdot|\mathbf{s}, \mathbf{a})} V(\mathbf{s}) \right). \quad (29)$$

We discuss the case with stochastic dynamics in the main paper which requires the estimation of \mathcal{B}^* . Under deterministic dynamic this can be avoided and we detail our method below.

A.3.1 \mathcal{X} -QL in Deterministic dynamics

For deterministic dynamics we can avoid having an expectation over the next states to simplify the Bellman equation.

We can develop two simple algorithms for this case without needing \mathcal{B}^* .

Value Iteration. We can write the value-iteration objective as:

$$Q(\mathbf{s}, \mathbf{a}) \leftarrow r(\mathbf{s}, \mathbf{a}) + \gamma V_\theta(\mathbf{s}'), \quad (30)$$

$$\mathcal{J}(\theta) = \mathbb{E}_{\mathbf{s} \sim \rho_\mu, \mathbf{a} \sim \mu(\cdot|\mathbf{s})} \left[e^{(Q(\mathbf{s}, \mathbf{a}) - V_\theta(\mathbf{s}))/\beta} - (Q(\mathbf{s}, \mathbf{a}) - V_\theta(\mathbf{s}))/\beta - 1 \right]. \quad (31)$$

Here, we learn a single model of the values $V_\theta(\mathbf{s})$ to directly solve Equation 29. For the current value estimate $V_\theta(\mathbf{s})$, we calculate targets $r(\mathbf{s}, \mathbf{a}) + \gamma V_\theta(\mathbf{s})$ and find a new estimate $V'_\theta(\mathbf{s})$ by fitting \mathbb{L}_μ^β with our objective \mathcal{J} . Using our Gumbel Regression framework, we can guarantee that as \mathcal{J} finds a consistent estimate of the \mathbb{L}_μ^β , and $V_\theta(\mathbf{s})$ will converge to the optimal $V(\mathbf{s})$ upto some sampling error.

616 **Q-Iteration.** Alternatively, we can develop a Q-iteration objective solving the recurrence:

$$Q_{t+1}(\mathbf{s}, \mathbf{a}) = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{L}_{\mathbf{a}' \sim \mu}^{\beta} [Q_t(\mathbf{s}', \mathbf{a}')] \quad (32)$$

$$= r(\mathbf{s}, \mathbf{a}) + \mathbb{L}_{\mathbf{a}' \sim \mu}^{\gamma\beta} [\gamma Q_t(\mathbf{s}', \mathbf{a}')] \quad (33)$$

$$= \mathbb{L}_{\mathbf{a}' \sim \mu}^{\gamma\beta} [r(\mathbf{s}, \mathbf{a}) + \gamma Q_t(\mathbf{s}', \mathbf{a}')] . \quad (34)$$

617 where we can rescale β to $\gamma\beta$ to move \mathbb{L} out.

618 This gives the objective:

$$Q^t(\mathbf{s}, \mathbf{a}) \leftarrow r(\mathbf{s}, \mathbf{a}) + \gamma Q_{\theta}(\mathbf{s}', \mathbf{a}'), \quad (35)$$

$$\mathcal{J}(Q_{\theta}) = \mathbb{E}_{\mu(\mathbf{s}, \mathbf{a}, \mathbf{s}')} \left[e^{(Q^t(\mathbf{s}, \mathbf{a}) - Q_{\theta}(\mathbf{s}, \mathbf{a}))/\gamma\beta} - (Q^t(\mathbf{s}, \mathbf{a}) - Q_{\theta}(\mathbf{s}, \mathbf{a}))/\gamma\beta - 1 \right]. \quad (36)$$

619 Thus, this gives a method to directly estimate Q_{θ} without learning values, and forms our \mathcal{X} -TD3
620 method in the main paper. Note, that β is a hyperparameter, so we can use an alternative hyperparam-
621 eter $\beta' = \gamma\beta$ to simplify the above.

622 We can formalize this as a Lemma in the deterministic case:

Lemma A.2. *Let*

$$\mathcal{J}(\mathcal{T}_{\mu}Q - Q') = \mathbb{E}_{\mathbf{s}, \mathbf{a}, \mathbf{s}', \mathbf{a}' \sim \mu} \left[e^{(\mathcal{T}_{\mu}Q(\mathbf{s}, \mathbf{a}) - Q'(\mathbf{s}, \mathbf{a}))/\gamma\beta} - (\mathcal{T}_{\mu}Q(\mathbf{s}, \mathbf{a}) - Q'(\mathbf{s}, \mathbf{a}))/\gamma\beta - 1 \right].$$

623 where \mathcal{T}_{μ} is a linear operator that maps Q from current (\mathbf{s}, \mathbf{a}) to the next $(\mathbf{s}', \mathbf{a}')$: $\mathcal{T}_{\mu}Q(\mathbf{s}, \mathbf{a}) :=$
624 $r(\mathbf{s}, \mathbf{a}) + \gamma Q(\mathbf{s}', \mathbf{a}')$

625 Then we have $\mathcal{B}^*Q^t = \underset{Q' \in \Omega}{\operatorname{argmin}} \mathcal{J}(\mathcal{T}_{\mu}Q^t - Q')$, where Ω is the space of Q -functions.

Proof. We use that in deterministic dynamics,

$$\mathbb{L}_{\mathbf{a}' \sim \mu}^{\gamma\beta} [\mathcal{T}_{\mu}Q(\mathbf{s}, \mathbf{a})] = r(\mathbf{s}, \mathbf{a}) + \gamma \mathbb{L}_{\mathbf{a}' \sim \mu}^{\beta} [Q(\mathbf{s}', \mathbf{a}')] = \mathcal{B}^*Q(\mathbf{s}, \mathbf{a})$$

626 Then solving for the unique minima for \mathcal{J} establishes the above results.

627 Thus, optimizing \mathcal{J} with a fixed-point is equivalent to Q-iteration with the Bellman operator.

628 □

629 B Gumbel Regression

630 We characterize the concentration bounds for Gumbel Regression in this section. First, we bound the
631 bias on applying \mathbb{L}^{β} to inputs containing errors. Second, we bound the error due to an empirical $\hat{\mathbb{L}}^{\beta}$
632 over finite N samples.

633 B.1 Overestimation Bias

634 Let $\hat{Q}(\mathbf{s}, \mathbf{a})$ be a random variable representing a Q-value estimate for a state and action pair (\mathbf{s}, \mathbf{a}) .
635 We assume that it is an unbiased estimate of the true Q-value $Q(\mathbf{s}, \mathbf{a})$ with $\mathbb{E}[\hat{Q}(\mathbf{s}, \mathbf{a})] = Q(\mathbf{s}, \mathbf{a})$. Let
636 $Q(\mathbf{s}, \mathbf{a}) \in [-Q_{max}, Q_{max}]$

637 Then, $V(\mathbf{s}) = \mathbb{L}_{\mathbf{a} \sim \mu}^{\beta} Q(\mathbf{s}, \mathbf{a})$ is the true value function, and $\hat{V}(\mathbf{s}) = \mathbb{L}_{\mathbf{a} \sim \mu}^{\beta} \hat{Q}(\mathbf{s}, \mathbf{a})$ is its estimate.

638 **Lemma B.1.** *We have $V(\mathbf{s}) \leq \mathbb{E}[\hat{V}(\mathbf{s})] \leq \mathbb{E}_{\mathbf{a} \sim \mu} [Q(\mathbf{s}, \mathbf{a})] + \beta \log \cosh(Q_{max}/\beta)$.*

639 *Proof.* The lower bound $V(\mathbf{s}) \leq \mathbb{E}[\hat{V}(\mathbf{s})]$ is easy to show using Jensen's Inequality as \log_sum_exp
640 is a convex function.

For the upper bound, we can use a reverse Jensen's inequality [34] that for any convex mapping f on the interval $[a, b]$ it holds that:

$$\sum_i p_i f(x_i) \leq f\left(\sum_i p_i x_i\right) + f(a) + f(b) - f\left(\frac{a+b}{2}\right)$$

Setting $f = -\log(\cdot)$ and $x_i = e^{\hat{Q}(\mathbf{s}, \mathbf{a})/\beta}$, we get:

$$\mathbb{E}_{\mathbf{a} \sim \mu}[-\log(e^{\hat{Q}(\mathbf{s}, \mathbf{a})/\beta})] \leq -\log(\mathbb{E}_{\mathbf{a} \sim \mu}[e^{\hat{Q}(\mathbf{s}, \mathbf{a})/\beta}]) - \log(e^{Q_{max}/\beta}) - \log(e^{-Q_{max}/\beta}) + \log\left(\frac{e^{Q_{max}/\beta} + e^{-Q_{max}/\beta}}{2}\right)$$

On simplifying,

$$\hat{V}(\mathbf{s}) = \beta \log(\mathbb{E}_{\mathbf{a} \sim \mu} e^{\hat{Q}(\mathbf{s}, \mathbf{a})/\beta}) \leq \mathbb{E}_{\mathbf{a} \sim \mu}[\hat{Q}(\mathbf{s}, \mathbf{a})] + \beta \log \cosh(Q_{max}/\beta)$$

Taking expectations on both sides, $\mathbb{E}[\hat{V}(\mathbf{s})] \leq \mathbb{E}_{\mathbf{a} \sim \mu}[Q(\mathbf{s}, \mathbf{a})] + \beta \log \cosh(Q_{max}/\beta)$. This gives an estimate of how much the LogSumExp overestimates compared to taking the expectation over actions for random variables \hat{Q} . This bias monotonically decreases with β , with $\beta = 0$ having a max bias of Q_{max} and for large β decaying as $\frac{1}{2\beta} Q_{max}^2$.

□

B.2 PAC learning Bounds for Gumbel Regression

Lemma B.2. $\exp(\hat{\mathbb{L}}^\beta(X)/\beta)$ over a finite N samples is an unbiased estimator for the partition function $Z^\beta = \mathbb{E}[e^{X/\beta}]$ and with a probability at least $1 - \delta$ it holds that:

$$\exp(\hat{\mathbb{L}}^\beta(X)/\beta) \leq Z^\beta + \sinh(X_{max}/\beta) \sqrt{\frac{2 \log(1/\delta)}{N}}.$$

Similarly, $\hat{\mathbb{L}}^\beta(X)$ over a finite N samples is a consistent estimator of $\mathbb{L}^\beta(X)$ and with a probability at least $1 - \delta$ it holds that:

$$\hat{\mathbb{L}}^\beta(X) \leq \mathbb{L}^\beta(X) + \frac{\beta \sinh(X_{max}/\beta)}{Z^\beta} \sqrt{\frac{2 \log(1/\delta)}{N}}.$$

Proof. To prove these concentration bounds, we consider random variables $e^{X_1/\beta}, \dots, e^{X_n/\beta}$ with $\beta > 0$, such that $a_i \leq X_i \leq b_i$ almost surely, i.e. $e^{a_i/\beta} \leq e^{X_i/\beta} \leq e^{b_i/\beta}$.

We consider the sum $S_n = \sum_{i=1}^N e^{X_i/\beta}$ and use Hoeffding's inequality, so that for all $t > 0$:

$$P(S_n - \mathbb{E}S_n \geq t) \leq \exp\left(\frac{-2t^2}{\sum_{i=1}^n (e^{b_i/\beta} - e^{a_i/\beta})^2}\right) \quad (37)$$

To simplify, we let $a_i = -X_{max}$ and $b_i = X_{max}$ for all i . We also rescale t as $t = Ns$, for $s > 0$. Then

$$P(S_n - \mathbb{E}S_n \geq Ns) \leq \exp\left(\frac{-Ns^2}{2 \sinh^2(X_{max}/\beta)}\right) \quad (38)$$

We can notice that L.H.S. is same as $P(\exp(\hat{\mathbb{L}}^\beta(X)/\beta) - \exp(\mathbb{L}^\beta(X)/\beta) \geq s)$, which is the required probability we want. Letting the R.H.S. have a value δ , we get

$$s = \sinh(X_{max}/\beta) \sqrt{\frac{2 \log(1/\delta)}{N}}$$

Thus, with a probability $1 - \delta$, it holds that:

$$\exp(\hat{\mathbb{L}}^\beta(X)/\beta) \leq \exp(\mathbb{L}^\beta(X)/\beta) + \sinh(X_{max}/\beta) \sqrt{\frac{2 \log(1/\delta)}{N}} \quad (39)$$

Thus, we get a concentration bound on $\exp(\hat{\mathbb{L}}^\beta(X)/\beta)$ which is an unbiased estimator of the partition function $Z^\beta = \exp(\mathbb{L}^\beta(X)/\beta)$. This bound becomes tighter with increasing β , and asymptotically behaves as $\frac{X_{max}}{\beta} \sqrt{\frac{2 \log(1/\delta)}{N}}$.

661 Similarly, to prove the bound on the log-partition function $\hat{\mathbb{L}}^\beta(X)$, we can further take $\log(\cdot)$ on both
 662 sides and use the inequality $\log(1+x) \leq x$, to get a direct concentration bound on $\hat{\mathbb{L}}^\beta(X)$,

$$\hat{\mathbb{L}}^\beta(X) \leq \mathbb{L}^\beta(X) + \beta \log \left(1 + \sinh(X_{max}/\beta) e^{-\mathbb{L}^\beta(X)/\beta} \sqrt{\frac{2 \log(1/\delta)}{N}} \right) \quad (40)$$

$$= \mathbb{L}^\beta(X) + \beta \sinh(X_{max}/\beta) e^{-\mathbb{L}^\beta(X)/\beta} \sqrt{\frac{2 \log(1/\delta)}{N}} \quad (41)$$

$$= \mathbb{L}^\beta(X) + \frac{\beta \sinh(X_{max}/\beta)}{Z^\beta} \sqrt{\frac{2 \log(1/\delta)}{N}} \quad (42)$$

663 This bound also becomes tighter with increasing β , and asymptotically behaves as $\frac{X_{max}}{Z^\beta} \sqrt{\frac{2 \log(1/\delta)}{N}}$.
 664 □

665 C Bridging soft and conservative Q-Learning

666 C.1 Inherent Conservatism in \mathcal{X} -QL

667 Our method is inherently conservative similar to CQL [19] in that it underestimates the value
 668 function (in vanilla Q-learning) $V^\pi(\mathbf{s})$ by $-\beta \mathbb{E}_{\mathbf{a} \sim \pi(\mathbf{a}|\mathbf{s})} \left[\log \frac{\pi(\mathbf{a}|\mathbf{s})}{\pi_{\mathcal{D}}(\mathbf{a}|\mathbf{s})} \right]$, whereas CQL underestimates
 669 values by a factor $-\beta \mathbb{E}_{\mathbf{a} \sim \pi(\mathbf{a}|\mathbf{s})} \left[\frac{\pi(\mathbf{a}|\mathbf{s})}{\pi_{\mathcal{D}}(\mathbf{a}|\mathbf{s})} - 1 \right]$, where $\pi_{\mathcal{D}}$ is the behavior policy. Notice that the
 670 underestimation factor transforms V^π in vanilla Q-learning into V^π used in the soft-Q learning
 671 formulation. Thus, we observe that KL-regularized Q-learning is inherently conservative, and this
 672 conservatism is built into our method.

673 Furthermore, it can be noted that CQL conservatism can be derived as adding a χ^2 regularization
 674 to an MDP and although not shown by the original work [19] or any follow-ups to our awareness,
 675 the last term of Eq. 14 in CQL's Appendix B [19], is simply $\chi^2(\pi || \pi_{\mathcal{D}})$ and what the original work
 676 refers to as D_{CQL} is actually the χ^2 divergence. Thus, it is possible to show that all the results for
 677 CQL hold for our method by simply replacing D_{CQL} with D_{KL} i.e. the χ^2 divergence with the KL
 678 divergence everywhere.

679 We show a simple proof below that D_{CQL} is the χ^2 divergence:

$$\begin{aligned} D_{CQL}(\pi, \pi_{\mathcal{D}})(\mathbf{s}) &:= \sum_{\mathbf{a}} \pi(\mathbf{a} | \mathbf{s}) \left[\frac{\pi(\mathbf{a} | \mathbf{s})}{\pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s})} - 1 \right] \\ &= \sum_{\mathbf{a}} (\pi(\mathbf{a} | \mathbf{s}) - \pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s}) + \pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s})) \left[\frac{\pi(\mathbf{a} | \mathbf{s})}{\pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s})} - 1 \right] \\ &= \sum_{\mathbf{a}} (\pi(\mathbf{a} | \mathbf{s}) - \pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s})) \left[\frac{\pi(\mathbf{a} | \mathbf{s}) - \pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s})}{\pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s})} \right] + \sum_{\mathbf{a}} \pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s}) \left[\frac{\pi(\mathbf{a} | \mathbf{s})}{\pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s})} - 1 \right] \\ &= \sum_{\mathbf{a}} \pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s}) \left[\frac{\pi(\mathbf{a} | \mathbf{s})}{\pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s})} - 1 \right]^2 + 0 \text{ since, } \sum_{\mathbf{a}} \pi(\mathbf{a} | \mathbf{s}) = \sum_{\mathbf{a}} \pi_{\mathcal{D}}(\mathbf{a} | \mathbf{s}) = 1 \\ &= \chi^2(\pi(\cdot | \mathbf{s}) || \pi_{\mathcal{D}}(\cdot | \mathbf{s})), \text{ using the definition of chi-square divergence} \end{aligned}$$

680 C.2 Why \mathcal{X} -QL is better than CQL for offline RL

681 In light of the above results, we know that CQL adds a χ^2 regularization to the policy π with respect
 682 to the behavior policy $\pi_{\mathcal{D}}$, whereas our method does the same using the reverse-KL divergence.

683 Now, the reverse-KL divergence has a mode-seeking behavior, and thus our method will find a policy
 684 that better fits the mode of the behavior policy and is more robust to random actions in the offline
 685 dataset. CQL does not have such a property and can be easily affected by noisy actions in the dataset.

686 C.3 Connection to Dual KL representation

For given distributions μ and π , we can write their KL-divergence using the dual representation proposed by IQ-Learn [12]:

$$D_{KL}(\pi || \mu) = \max_{x \in \mathbb{R}} \mathbb{E}_{\mu}[-e^{-x}] - \mathbb{E}_{\pi}[x] - 1,$$

687 which is maximized for $x = -\log(\pi/\mu)$.

688 We can make a clever substitution to exploit the above relationship. Let $x = (Q - \mathcal{T}^{\pi} \hat{Q}^k)/\beta$ for a
689 variable $Q \in \mathbb{R}$ and a fixed constant $\mathcal{T}^{\pi} \hat{Q}^k$, then on variable substitution we get the equation:

$$\mathbb{E}_{s \sim \rho_{\mu}}[D_{KL}(\pi(\cdot|s) || \mu(\cdot|s))] = \min_Q \mathcal{L}(Q), \text{ with}$$

690

$$\mathcal{L}(Q) = \mathbb{E}_{s \sim \rho_{\mu}, a \sim \mu(\cdot|s)} \left[e^{(\mathcal{T}^{\pi} \hat{Q}^k(s, a) - Q(s, a))/\beta} \right] - \mathbb{E}_{s \sim \rho_{\mu}, a \sim \pi(\cdot|s)} [(\mathcal{T}^{\pi} \hat{Q}^k(s, a) - Q(s, a))/\beta] - 1$$

691 This gives us Equation 8 in Section 3.3 of the main paper, and is minimized for $Q = \mathcal{T}^{\pi} \hat{Q}^k -$
692 $\beta \log(\pi/\mu)$ as we desire. Thus, this lets us transform the regular Bellman update into the soft-Bellman
693 update.

694 D Broader Impacts

695 We introduce new methods for offline and online RL. RL has the potential to power modern robotics
696 and decision making systems. Care must be taken, however, to ensure that the objectives of these
697 systems are well aligned with humans so that they can improve lives.

698 E Experiments

699 In this section we provide additional experimental results and further details on all experimental
700 procedures.

701 E.1 \mathcal{X} -QL offline Learning Curves




Figure 4: **Offline RL Results.** We show the returns vs number of training iterations. (Averaged over 5 seeds)

702 E.2 Additional Offline Environments

Table 3: Evaluation on Franka Kitchen and Adroit tasks from D4RL

Dataset	BC	BRAC-p	BEAR	Onestep RL	CQL	IQL	\mathcal{X} -QL
pen-human-v0	63.9	8.1	-1.0	-	37.5	71.5	91.4
hammer-human-v0	1.2	0.3	0.3	-	4.4	1.4	5.4
door-human-v0	2	-0.3	-0.3	-	9.9	4.3	10.1
relocate-human-v0	0.1	-0.3	-0.3	-	0.2	0.1	1.1
pen-cloned-v0	37	1.6	26.5	60.0	39.2	37.3	58.2
hammer-cloned-v0	0.6	0.3	0.3	2.1	2.1	2.1	4.6
door-cloned-v0	0.0	-0.1	-0.1	0.4	0.4	1.6	8.9
relocate-cloned-v0	-0.3	-0.3	-0.3	-0.1	-0.1	-0.2	-0.1

703 E.3 Additional Bellman Error Plots




Figure 5: Additional plots of the error distributions for different environments. We find that the Gumbel distribution strongly fit the errors in first two environments, Cheetah and Walker, but provides a worse fit in the Hopper environment. Nonetheless, we see performance improvements in Hopper using our approach.

704 E.4 Numeric Stability

705 In practice, a naive implementation of the Gumbel loss function \mathcal{J} from Equation 11 suffers from
 706 stability issues due to the exponential term. We found that stabilizing the loss objective was essential
 707 for training. Practically, we follow the common max-normalization trick used in softmax computation.
 708 This amounts to factoring out $e^{\max_z z}$ from the loss and consequently scaling the gradients. We
 709 additionally clip loss inputs that are too large. An example code snippet in Pytorch is included below:

```
def gumbel_loss(pred, label, beta, clip):
    z = (label - pred)/beta
    z = torch.clamp(z, -clip, clip)
    max_z = torch.max(z)
    max_z = torch.where(max_z < -1.0, torch.tensor(-1.0), max_z)
    max_z = max_z.detach() # Detach the gradients
    loss = torch.exp(z - max_z) - z*torch.exp(-max_z) - torch.exp(-max_z)
    return loss.mean()
```

710 In some experiments we additionally clip the value of the gradients for stability.

711 E.5 Offline Hyperparameters

712 We base our implementation of \mathcal{X} -QL off of the official implementation of IQL from Kostrikov et al.
 713 [17], and use their code to run our experiments. We use the same network architecture and also apply
 714 the Double- Q trick. We also apply the same data preprocessing which is described in their appendix.
 715 We additionally take their baseline results and use them in Table 1, Table 2, and Table 3 for accurate
 716 comparison. Kostrikov et al. [17] did not release code for Adroit Binary-v0, and as such we do not
 717 have comparison on these tasks. We contacted the authors regarding the release of these environments
 718 and are awaiting a response.

719 We keep our general algorithm hyper-parameters and evaluation procedure the same but tune β and
 720 the gradient clipping value for each environment. For MuJoCo locomotion tasks we average mean
 721 returns over 10 evaluation trajectories and 5 random seeds. For the AntMaze tasks, we average over
 722 1000 evaluation trajectories. When reporting results for the Adroit environments we use smoothing
 723 of 0.7 as evaluations fluctuate significantly.

724 E.6 Online Hyperparameters

725 We base our implementation of SAC off pytorch_sac [40] but modify it to use a Value function
 726 as described in Haarnoja et al. [13]. Empirically we see similar performance with and without
 727 using the value function, but leave it in for fair comparison against our \mathcal{X} -SAC variant. We base
 728 our implementation of TD3 on the original author’s code from [10]. Like in offline experiments,
 729 hyper-parameters were left as default except for β , which we tuned for each environment. In online
 730 experiments we clipped the exponential term to be at maximum 8. For SAC we ran three seeds
 731 in each environment as it tended to be more stable. For TD3 we ran four. Occasionally, our \mathcal{X} -

732 variants would experience exploding loss terms causing performance to collapse. We re-ran these
733 seeds to obtain runs that did not become numerically unstable. For all online experiments we used an
734 exponential clip value of 8. We verified our SAC results by comparing to [40] and our TD3 results by
735 comparing to [21]. We found that our TD3 implementation performed marginally better overall.

Table 4: Offline RL Hyperparameters used for \mathcal{X} -QL

Env	Beta	Grad Clip
halfcheetah-medium-v2	1	7
hopper-medium-v2	5	7
walker2d-medium-v2	10	7
halfcheetah-medium-replay-v2	1	7
hopper-medium-replay-v2	2	7
walker2d-medium-replay-v2	5	7
halfcheetah-medium-expert-v2	1	7
hopper-medium-expert-v2	1	7
walker2d-medium-expert-v2	2	7
antmaze-umaze-v0	2	7
antmaze-umaze-diverse-v0	5	7
antmaze-medium-play-v0	0.8	5
antmaze-medium-diverse-v0	1	5
antmaze-large-play-v0	0.8	5
antmaze-large-diverse-v0	0.5	5
kitchen-complete-v0	2	7
kitchen-partial-v0	5	7
kitchen-mixed-v0	2	7

Table 5: Hyperparameters for online RL Algorithms

Parameter	SAC	TD3
Batch Size	1024	256
Learning Rate	0.0001	0.001
Critic Freq	1	1
Actor Freq	1	2
Actor and Critic Arch	1024, 1024	256, 256
Buffer Size	1,000,000	1,000,000
Actor Noise	Auto-tuned	0.1, 0.05 (Hopper)
Target Noise	–	0.2

Table 6: Values of β used for online experiments

Env	\mathcal{X} -SAC	\mathcal{X} -TD3	\mathcal{X} -TD3 - DQ
Cheetah Run	2	5	4
Walker Run	1	2	4
Hopper Hop	2	2	3
Quadruped Run	5	5	20

736 E.7 Error Figure 1

Figure 1 was generated by running the SAC implementation from Yarats and Kostrikov [40] for 100,000 timesteps on the Cheetah Run environment and the logging the Bellman errors every 5,000 steps. In particular, the Bellman errors were computed as:

$$r(\mathbf{s}, \mathbf{a}) + \gamma Q_{\theta_1}(\mathbf{s}', \pi_{\psi}(\mathbf{s}')) - Q_{\theta_1}(\mathbf{s}, \mathbf{a})$$

737 In the above equation Q_{θ_1} represents the first of the two Q networks used in the Double Q trick.
738 We do not use target networks to compute the bellman error, and instead compute the fully online
739 quantity. $\pi_{\psi}(\mathbf{s}')$ represents the mean or deterministic output of the current policy distribution. Also
740 note that the entropy term was not added as we seek to characterize the standard bellman error and
741 not the soft-bellman error. Before generating the plot the errors were clipped to the ranges shown in
742 the plot. This tended prevented over-fitting to large outliers.

743 E.8 Resources

744 We ran experiments largely on Nvidia Titan XP GPUs via an internal cluster. Tuning values of β was
745 done largely via trial and error.